

# Detecting TravisCI Misconfiguration

## Plan of Attack

Name1  
Affiliation1  
Email1

Name2    Name3  
Affiliation2/3  
Email2/3

### Abstract

We can detect TravisCI misconfigurations using ConfigC.

### 1. Availability of Data

TravisCI has recently released a metadata api to study their tool[?].

The first use of this API was analyzing the metadata of builds. It was found that ~15-20% of failed TravisCI builds are due to "errors". My understanding of the API is that this corresponds to a misconfiguration. This is a large enough number that not only will reducing this be more convenient for users, but it could also significantly reduce server costs for Travis. By using the `tr_duration` field from the API, we can also figure out how much server time we can save[?].

### 2. Feasability Analysis

The language of TravisCI configuration files can likely be handled by ConfigC. ConfigC works on languages with shallow parse trees, and most of the `.travis.yml` files is shallow. However, because `.travis.yml` allows for bash scripts, some error may be out of scope. We want to identify the frequency of the types of errors - if too many are related to poorly formed bash scripts, a naive application of ConfigC will not suffice.

To confirm that the types of most misconfigurations are in the scope of ConfigC, we should compile a collection of common misconfigurations. Usually such a task requires a domain expert. However, in this case we may be able to learn the common errors by using the API and the temporal data from the commits.

We might code an error as the diff between a sequential pair of an erroring commit and a passing commit. As a new direction, we may also explore how useful this data is for learning/verification directly. My guess is that is data will be too noisy, but either way, it should give us a sense of the problem.

To build such a database, scrape the `.travis.yml` files from each commit for a set of repos. This can be achieved with the `tr_status`, `git_commit` and `gh_project_name` fields from the API.

### 2.1 Learning from Temporal properties

An interesting extension to ConfigC on a theoretical side is to look at the temporal relations between commits. The key observation is that with each commit that changes the build status, we can learn highly localized information about our model.

The first step is to build an intermediate representation of the data we will learn. For the sake of efficiency, this data must be structured as a shallow tree. While the `travis.yml` file can be taken verbatim, we must use a simplification of the program code. Since we are only interested in the parts of the program effecting build status, we should extract the imported libraries as a list. We will call this summary  $S_t$ , a simplified representation of the program which contains only the information relevant to the learning process. The subscript is a tag based on the ordered commit history. To handle branches, add a superscript to indicate the branch, and restart the counter on a branch - merging should be handled easily (not exactly sure how yet).

From this summary we can then build a model  $M(P_t)$ , which is the full set of possible rules derivable from the model. These possible rules must be given by the user. In [?], these were called Attributes, and express constraints such as, ordering of lines, or substring relations on values.

We will denote the build status of  $S_t$  with  $B(S_t)$ . In this application, we consider only the passing and erroring build status, denoted  $P$ , and  $E$  respectively. All status that are not `error`, as defined by the Travis API, will be included as passing. For brevity, we denote sequences of build statuses with the following notation:

$$S(P_t) = P \wedge S(P_{t+1}) = E \implies S(P_{t,t+1}) = PE$$

We know that if a build is passing, then there must not exist any errors. That is, the model must not contain any rules which are breaking. Note we are not, however, guaranteed that any rules from a passing commit are necessary.

$$S(P_t) = P \implies \forall r \in M(P_t), r \notin Br$$

The key insight is that when we commit a break, we can localize the error to one of the lines that changed. Either we removed something that was necessary, or added something that was breaking. We use an inclusive disjunction, since a erroring commit can break multiple things at once. Expressed formally, where  $\setminus$  is the set difference), that is:

$$\begin{aligned} S(P_{t,t+1}) = PE &\implies \\ \exists r \in (M(P_t) \setminus M(P_{t+1})), r \in Nec \vee \\ \exists r \in (M(P_{t+1}) \setminus M(P_t)), r \in Br \end{aligned}$$

We then can combine all these formulas with conjunctions and send it to an SMT solver and magically get an answer. While existential set operations can be expensive in general for an SMT solver, in our application this is not the case. Because these sets will be small thanks to the practice of making incremental commits

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

CONF 'yy,    Month d-d, 20yy, City, ST, Country.  
Copyright © 20yy ACM 978-1-xxxx-xxxx-n/yy/mm...\$15.00.  
<http://dx.doi.org/10.1145/nnnnnnnn.nnnnnnn>

when using source control, these operations will be fairly cheap. In fact, the above implication generalizes to  $P_{t,t+n}$ , but for efficiency we must require that  $M(P_t) \setminus M(P_{t+n})$  is managably small. The definition of small here remains to be experimentally determined.

### 3. Implementation

To build the learning set, we will use the API to build tuples of `.travis.yml` and repo info. In addition to the `.travis.yml` file (collected as above), we will need at a minimum `gh_lang` and `tr_status`. With this `(File,gh_lang,tr_status,...)` tuple, we can then (almost) directly apply ConfigC to detect errors in the shallow part of the tree.

The first (and simplest) requirement is a `.travis.yml` parser, which is an extension of normal yml parsing to handle travis' ability to include bash scripts. Second, depending on the results from Sec 2, we may also require some extra domain knowledge of the types of errors to be encoded as possible rules. This is a slightly more involved task that requires writing some non-trivial Haskell code - but it really isn't that bad.